

Learning to Extract Events without Human-annotated Text

Abstract

Existing event extraction systems are often supervised and rely on expert-annotated datasets, with limited event types. However, designing and constructing these high-quality corpora, usually with limited size and coverage of event types, is costly, which makes learned extractors hard to generalize. With the essence of distant supervision, we investigate the possibilities of automatic construction of training data for various event types with the help of structured knowledge bases. We further propose a novel neural network with ILP-based post inference committing to handling two challenges in event extraction: multi-type events and multi-word arguments. Both automatic and manual evaluations demonstrate that it is possible to learn to extract various events, according to existing knowledge bases, without human-annotated training data.

1 Introduction

Automatically extracting events from natural text remains a challenging task in information extraction. Among diverse types of event extraction systems, the extraction task proposed by Automatic Content Extraction (ACE) [Doddington *et al.*, 2004] is the most popular framework, which defines two main terminologies: **trigger** and **argument**. The former is the word that most clearly expresses the occurrence of an event. The latter is a phrase that serves as a participant or attribute with a specific role in an event.

However, constructing training data for ACE task is expensive. First, linguists are required to summarize a large amount of text to elaborately design templates about potential arguments for each event type. Second, rules should be explicitly stated to guide annotators. In spite of detailed guidelines, there is still disagreement among human annotators about what should (not) be regarded as triggers/arguments. For example, can a prepositional phrase or a portion of a word trigger an event, e.g., *in prison* triggers an *arrest* event, or, *ex in ex-husband* triggers a *divorce* event? Besides, ACE event extraction systems remain two major limitations: single-token trigger labeling, and one type for one event.

It would be interesting to see (1) can we automatically build a dataset for event extraction without experts involved?

and (2) can we have an event extractor that handles more realistic scenarios, e.g., when trigger annotations are unavailable, or events with more than one type.

First, we observe that structured knowledge bases (KB) often organize complex structured information in tables, which share similar structures with ACE event definitions. A particular entry of such tables usually implies the occurrence of certain events. On the other hand, recent studies [Mintz *et al.*, 2009; Zeng *et al.*, 2015] have demonstrated the effectiveness of KB as distant supervision for binary relation extraction. However, there are two major challenges when leveraging KB to event extraction: first, event structures are more complex than binary relations. They can be represented as $\langle event_type, argument_1, \dots, argument_n \rangle$, which are n -ary relations with various numbers of arguments. Second, there is no explicit trigger information in any existing knowledge base. Therefore, to explore the distant supervision assumption in event extraction, we investigate different hypotheses for better data quality and quantity. Among them, the vital one is that, for a particular event type, there is a group of **key arguments** which together can imply an event instead of explicit triggers. We utilize Freebase as our knowledge base and Wikipedia articles as text for data generation. According to Mintz *et al.* [2009], because a major source of Freebase is the tabular data from Wikipedia, making it a natural fit with Freebase. Figure 1 illustrates examples of sentences annotated by our algorithm.

Second, unlike previous studies focus on tasks defined by ACE evaluation framework [Ahn, 2006; Li *et al.*, 2013; Chen *et al.*, 2015; Nguyen *et al.*, 2016], we propose a novel event extraction paradigm with key arguments to characterize an event type. We consider event extraction as two sequence labeling subtasks, namely event detection and argument detection. Inspired by neural network models in sequence labeling tasks [Huang *et al.*, 2015; Lample *et al.*, 2016], we utilize LSTM-CRF models to label key arguments and non-key arguments in the each sentence separately. However, event structures are not simple sequences and there are strong dependencies among key arguments. We therefore reformulate the hypotheses as constraints, and apply linear integer programming to output multiple optimal label sequences to capture multi-type events.

In this paper, we exploit existing structured knowledge bases, e.g., Freebase, as distant supervision to automati-

cally annotate event structures from plain text without human annotations. We further propose a novel event extraction paradigm that harnesses key arguments to imply certain event types without explicit trigger annotations. We present an LSTM-CRF model with post inference to extract both Freebase-style events as well as multi-type event mentions on the generated dataset, which is demonstrated effective by both manual and automatic evaluations.

2 Dataset Preparation

We employ Freebase to automatically annotate text in Wikipedia. We regard a sentence as a positive one when it mentions an occurrence of event, or otherwise a negative sentence. For example, S1 and S2 are positive sentences and their arguments are in *italics* and underlined, while S3 and S4 are negative sentences. The event structures of S1 and S2 are illustrated in Figure 1.

S1: *Remedy Corp* was sold to *BMC Software* as the *Service Management Business Unit* in 2004.

S2: *Microsoft* spent \$6.3 billion buying online display advertising company *aQuantive* in 2007.

S3: Microsoft hopes aQuantive’s Brian McAndrews can outfox Google.

S4: On April 29th, Elizabeth II and Prince Philip witnessed the marriage of Prince William.

2.1 Freebase

Freebase[Bollacker *et al.*, 2008] is a collaborative structured knowledge base which is divided into three layers: *domain*, *type* and *instance*. *Instances* are entries in Freebase, and are related to real-world entities like people and places. *Types* are different perspectives of *instances*. **Compound Value Type** (CVT) is a special type in Freebase to represent complex structured data where each instance consists of multiple *properties*. Some of the CVT schemas are highly alike to event structures where CVT properties can be treat as event arguments. For example, in Figure 1, *business.acquisition* is a CVT whose properties are *company_acquired*, *acquiring_company*, *date* and *divisions_formed*. These properties can also be used to represent participants and attributes in events extracted from S1 and S2.

We use the Freebase version of Berant *et al.* [2013], containing 1010 CVTs. After filtering out CVTs that describe the structures of the Freebase or are irrelevant to event extraction (e.g., *food.recipe_ingredient*), we select 24 CVTs with around 280 million instances.

2.2 Data Generation

H1: Positive sentences should contain all properties

This hypothesis indicates that if a sentence has all properties of a CVT, it is more likely to be a positive sentence. We regard the CVT as event type and extract words and phrases that match the properties of a CVT instance as involved arguments. For example, S1 contains all the properties of instance *m.07bh4j7* whose type is *business.acquisition*, thus we consider S1 as a positive sentence which implies an event about

Instances of <i>business.acquisition</i> in Freebase				
id	property	company_acquired	acquiring_company	date
m.07bh4j7		Remedy Corp	BMC Software	2004
m.05nb3y7		aQuantive	Microsoft	2007
				divisions_formed
				Service Management Business Unit
				NONE

↓ Data generation

Event structures in our dataset				
Wiki text	S1: <i>Remedy Corp</i> was sold to <i>BMC Software</i> as the <i>Service Management Business Unit</i> in <u>2004</u> .			
Event type	<i>business.acquisition</i>			
Arguments	company_acquired	acquiring_company	date	divisions_formed
	<i>Remedy Corp</i>	<i>BMC Software</i>	<u>2004</u>	<i>Service Management Business Unit</i>

Wiki text	S2: <i>Microsoft</i> spent \$6.3 billion buying online display advertising company <i>aQuantive</i> in <u>2007</u> .			
Event type	<i>business.acquisition</i>			
Arguments	acquiring_company	company_acquired	date	divisions_formed
	<i>Microsoft</i>	<i>aQuantive</i>	<u>2007</u>	—

Figure 1: Examples of CVT instances in Freebase, and labeled sentences in our dataset. *Company_acquired*, *acquiring_company* and *date* are key arguments in *business.acquisition*.

CVT	Key properties
award.award_honor	award_winner, award, . . . , year
film.performance	actor, film, character
education.education	institution, student, end_date
business.employment_tenure	company, title, person, from

Table 1: Examples of key properties of four CVTs.

business.acquisition, and *BMC Software*, *Remedy Corp*, *Service Management Business Unit* and *2004* should be labeled as the arguments that play the role of *acquiring_company*, *company_acquired*, *divisions_formed*, *date*, respectively.

However, in practice, we realize that *H1* is too strict that excludes a great many positive sentences like S2.

H2: Positive sentences should contain all key properties

This hypothesis is an extension of *H1*, which relaxes “all properties” constraint to “key properties”. We define the CVT property that plays an important part in its CVT structure and helps to distinguish with other CVT as **key property**. And **key argument** is the word/phrase that matches a key property of a CVT instance. For example, *company_acquired* and *acquiring_company* are the key properties of CVT *business.acquisition*, and with this relaxation, positive sentences like S2 need not contain all properties, but only key properties instead. Table 1 lists key properties of four CVTs.

The importance of a property *prop* (e.g., *date*) to its CVT *cvt* (e.g., *business.acquisition*) can be defined as follows:

$$degree_{cvt,prop} = \log \frac{count(cvt, prop)}{count(cvt) \times count(prop)} \quad (1)$$

where $count(cvt)$ is the number of all *cvt* instances, $count(prop)$ is the number of *prop* among all CVTs, and $count(cvt, prop)$ is the number of *cvt* instances that contain the property *prop*.

H3: Key properties should include time property

We discover that for many CVTs, their key properties do not take into account time property. However, ignoring time property will produce a large number of negative sentences like S3. It does not express an explicit event about




Figure 2: An illustration of dependency parse tree of S4.

business.acquisition while contain all key properties of an instance, resulting in mistaking *Microsoft* for *acquiring_company*, and *aQuantive* for *company_acquired*. By adding *date* to the set of key properties, S3 will be filtered. Therefore, we choose the time property which achieves the highest importance value as a supplementary key property.

H4: Positive sentences should contain key properties with close syntactic distance

We introduce another factor, syntactic distance, to annotate positive sentences. Intuitively, two arguments participant in the same event are likely to have close syntactic distance. This factor is effective to eliminate negative sentences, such as S4. The syntactic distance can be measured by the distance of two words in dependency parsing tree. We set the maximum distance between two key arguments as 2, denoting that, for a candidate sentence, if a pair of key arguments violates this constraint, it is supposed to be negative. Given the dependency parsing tree in Figure 2, S4 is negative because the distance between *Prince Philip* and *marriage* is 3.

We conduct a manual evaluation on the quantity and quality of datasets generated by different hypotheses (see Section 4.2), and utilize the combinations of hypothesis *H3* and *H4* as the final strategy.

3 Model

Unlike existing event extraction work in which triggers are the key clues to identify event and classify different event types, in the absence of human-labeled triggers, we argue that **key arguments** can play the same role as triggers. Consequently, we can treat event extraction as a pipeline of two primary subtasks, namely event detection and argument detection.

Event detection aims to identify key arguments in the sentence. And if a sentence contains **all** key arguments of a specific event type, it is considered to imply an event of the corresponding type. Take S1 as an example, *Remedy Corp*, *BMC Software*, and *2004* should be identified as *company_acquired*, *acquiring_company*, and *date*, respectively. As a result, S1 should be labeled as expressing an event about *business.acquisition*.

Argument detection aims to identify other non-key arguments for each event in the sentence. For the *business.acquisition* event in S1, *Service Management Business Unit* should be identified as *divisions_formed*.

3.1 Event Detection

Before presenting our model, we need to propose the solution to multi-words arguments. Then we introduce the components in our LSTM-CRF-ILP_{multi} model one-by-one from bottom to top.

Tagging scheme 68 percent of the arguments in our dataset consist of more than one words. To address this issue, we model each subtask as a sequence labeling task rather than a word classification task. Each word in the given sentence is tagged in the BIO scheme, where each token is labeled as B-role if it is the beginning of an event argument with its corresponding role, or I-role if it is inside an argument, or O otherwise.

LSTM Long Short-Term Memory Network (LSTM) [Hochreiter and Schmidhuber, 1997] is a natural fit for sequence labeling task, which maintains a memory based on historical contextual information. Formally, given a sentence $w = \{w_1, w_2, \dots, w_n\}$ of length n , we use \mathbf{x}_t to represent feature vector (e.g. word embedding) corresponding to the t -th word w_t . At each time step t , a LSTM unit takes \mathbf{x}_t as input and computes the output vector \mathbf{h}_t through several multiplicative gates. Then this output vector is fed into a softmax layer to estimate a probability distribution over all possible labels.

CRF A straightforward way is to choose the label which obtains maximum probability by LSTM as the prediction for each word. However, this independent labeling strategy is limited especially when there are strong dependencies and constraints between labels. To model the correlations between labels, we introduce a CRF layer into the output of LSTM, which is widely used and effective in various sequence labeling tasks, such as POS tagging and NER [Collobert *et al.*, 2011; Huang *et al.*, 2015].

We consider \mathbf{P} to be a matrix of confidence scores output by LSTM network, and the element $\mathbf{P}_{i,j}$ of the matrix denotes the probability of the label j for the i -th word in a sentence. The CRF layer has a transition score matrix \mathbf{A} as parameter, where $\mathbf{A}_{i,j}$ represents the score of a transition from label i to label j . The score of a sentence w along with a path of labels $y = \{y_1, y_2, \dots, y_n\}$ is measured by the sum of neural network outputs and transition scores:

$$\text{score}(w, y) = \sum_{i=0}^n \mathbf{P}_{i,y_i} + \sum_{i=1}^n \mathbf{A}_{y_i, y_{i+1}}, \quad (2)$$

During test, given a sentence w , we adopt the Viterbi algorithm [Rabiner, 1989] to find the optimal label sequence with the maximum score among all possible label sequences.

ILP-based Post Inference Event detection is a structure prediction problem, while the output sequences of LSTM-CRF not necessarily satisfy the structural constraints. Specifically, regardless of how many key arguments are identified correctly by LSTM-CRF, if there is one key argument missing, detection of its corresponding event is failed. To amend this flaw, we apply Integer Linear Programming (ILP) with respect to the scores given by the above LSTM-CRF model to generate the final labeling sequence.

Formally, let \mathcal{L} be the set of possible argument labels. For each word w_i in the sentence w and a pair of labels $\langle l, l' \rangle \in \mathcal{L} \times \mathcal{L}$, we create a binary variable $v_{i,l,l'} \in \{0, 1\}$, denoting

whether or not the i -th word w_i is tagged as label l and its following word w_{i+1} is tagged as label l' at the same time. The objective of ILP is to maximize the overall score of the variables,

$$\sum_{i,l,l'} v_{i,l,l'} * (\mathbf{P}_{i,l} + \mathbf{A}_{l,l'}).$$

We consider the following four constraints:

C1: Each word should be and only be labeled with one label, i.e.:

$$\sum_{l,l'} v_{i,l,l'} = 1 \quad (3)$$

C2: If the value of $v_{i,l,l'}$ is 1, then there has to be a label l^* which makes v_{i+1,l',l^*} equal to 1, i.e.:

$$v_{i,l,l'} = \sum_{l^*} v_{i+1,l',l^*} \quad (4)$$

C3: If the current label is I-arg, then its previous label must be B-arg, i.e.:

$$v_{i,\text{I-arg},l'} = v_{i-1,\text{B-arg},\text{I-arg}} \quad (5)$$

C4: For a specific event type, its key arguments should co-occur, or none of them should appear in the resulting sequence. For any pair of key arguments arg_1 and arg_2 with respect to the same event type, the variables related to them are subject to:

$$\sum_{i,l'} v_{i,\text{B-arg}_1,l'} \leq n * \sum_{j,l^*} v_{j,\text{B-arg}_2,l^*} \quad (6)$$

where n is the length of the sentence.

In order to address the multi-type event mention issue, we allow ILP solver to output multiple sequences iteratively. Formally, let $\mathbf{s}^t = \{l_1^t, l_2^t, \dots, l_n^t\}$ be the sequence produced at iteration t . During each iteration t , we eliminate $\{\mathbf{s}^1, \dots, \mathbf{s}^{t-1}\}$ from the solution space to obtain the next optimal sequences \mathbf{s}^t . We repeat the above procedure with these constraints through ILP, until the difference between objective value of \mathbf{s}^1 and \mathbf{s}^T is greater than a threshold λ , and consider all sequences $\{\mathbf{s}^1, \mathbf{s}^2, \dots, \mathbf{s}^{T-1}\}$ as the optimized set of label sequences. In our experiment, Gurobi [Gurobi Optimization, 2016] is chosen as our ILP problem solver and $\lambda = 0.05 \times n$.

3.2 Argument Detection

After event detection, a sentence will be classified into different event types, and labeled with its corresponding key arguments. Next step is argument detection which aims to identify the remaining arguments (non-key arguments) in the sentences.

We simply adapt the same architecture as the LSTM-CRF model (see Section 3.1) for argument detection, where we encode the label (output of event detection) of each word into a key-argument feature vector through a look-up table, and concatenate it with the origin word embedding as the input vector to LSTM. We do not need post inference in this task.

4 Experiments

4.1 Experimental Setup

Dataset and Evaluation Methodology. We use the November 20th, 2016 English Wikipedia dump, and generate 7180

sentences, containing 7376 events and 25840 arguments as dataset. We then randomly select 4800 sentences for training and 1180 sentences as test set, and the rest 1200 sentences for validation. We conduct both automatic evaluation and manual evaluation in the experiments. We first manually evaluate the quality of our test set. Next, we regard the noisy generated data as gold standard and evaluate our model automatically. Finally, we manually evaluate a subset of events detected by our model and analysis the differences with results in automatic evaluation.

Evaluation Measures. We evaluated our models in terms of precision (P), recall (R), and F-measure (F) for each subtask. These performance metrics are computed according to the following standards of correctness: For event type classification, an event is correctly classified if its reference sentence contains all key arguments of this event type; For argument detection, an argument is correctly detected if its offsets, role, and related event type exactly match any reference argument within the same sentence; For event detection, an event is correctly detected if its type and all its key arguments match a reference event within the same sentence.

Training. All hyperparameters are tuned on the development set. In event detection, we set the size of word embedding to 200, the size of LSTM layer to 100. In argument detection, we use the same size of word embedding, while the size of LSTM layer is 150, and the size of key argument embedding is 50. Word embeddings are pretrained using skip-gram word2vec model [Mikolov *et al.*, 2013] over the whole Wikipedia dump and fine tuned during training. To mitigate overfitting, we apply a dropout rate of 0.5 on both the input and output layers.

4.2 Dataset Evaluation

For comparison, we evaluate five datasets that utilize different hypotheses to generate positive sentences from Wikipedia. We randomly select 100 sentences in each dataset, and annotators are asked to determine whether these sentences imply events.

Hypothesis	H1	H2	H2+H4	H3	H3+H4
Instances	0.3M	3.6M	3.6M	1.3M	1.3M
Dataset	203	108K	12K	9241	7180
Event type	9	24	24	24	24
Correct (%)	98	22	37	81	89

Table 2: Statistic of generated dataset with different hypotheses. Instances denotes the number of CVT instances that can be used for each hypothesis. Dataset is the number of generated sentences. Event type indicates the number of different CVT types in each dataset. Correct represents the percentage of sentences which account as stating events explicitly.

As shown in Table 2, the strictest hypothesis, *H1*, guarantees the quality and confidence of generated data, while we can merely obtain 203 sentences and cover 9 types of events, which is quite insufficient for further training. *H2* is looser than *H1*, though expands the resulting dataset, it produces a large number of noisy sentences. This side effect demonstrates that *H2* is inappropriate to be used as a soft constraint.

Compared with *H2*, the significant improvement in the quality of sentences generated by *H3* proves that CVT properties referring time information are critical to data generation. Among all hypotheses, finally, data obtained by a combination of *H3* and *H4* achieves highest precision, which demonstrates that our hypothesis *H3* and *H4* are feasible and it is an effective way to generate reliable data automatically.

4.3 Baselines

To investigate the effectiveness of our proposed model, we develop three baseline extraction systems for comparison, including traditional feature-based methods and neural network models. For neural network method, we train a long short-term memory network that takes word embeddings as the input, and simply learns a probability distribution over all possible labels. For feature-based methods, we apply Conditional Random Field [Lafferty *et al.*, 2001] and Maximum Entropy [Berger *et al.*, 1996] to explore a variety of elaborate features (lexical, syntactic, entity information features) modified from state-of-art feature-based ACE event extraction system [Li *et al.*, 2013]. It is worth mentioning that during argument detection, we add the label of each word output by event detection as a supplementary feature.

We derive these features using Stanford CoreNLP [Manning *et al.*, 2014], and apply the implementation from the CRF++ toolkit [Kudo, 2005] and Le Zhang¹ to train CRF and max entropy classifiers, respectively.

4.4 Automatic Evaluations

As we can summarize from Table 3, traditional feature-based models are inefficient in both event detection and argument detection. Some of features they utilized, such as dependency features, suffer much from the absence of trigger. Although they can achieve high precisions, they can only extract a limited number of events, resulting in low recalls. Neural-network-based methods performs much better than feature-based models, because they can make better use of word semantic features, especially, LSTM can capture longer range dependencies and richer contextual information instead of neighbor word features.

Effect of CRF Layer. Every model which has a CRF layer over its LSTM output layer is superior to the one with a simple LSTM layer. Compared with LSTM model, LSTM-CRF achieves higher precisions and recalls in all subtasks by significantly reducing the invalid labeling sequences (e.g., *I-arg* appears right after *O*). During prediction, instead of tagging each token independently, LSTM-CRF takes into account the constraints between neighbor labels, and increases the cooccurrences of key arguments with regard to the same event type in some way.

Effect of Post Inference. As shown in Table 3, post inference based on ILP considerably improve the overall system performance, especially in event classification. With the help of constraint **C4**, some dubious key arguments can be inferred through other key arguments from their contexts. Compared

with LSTM-CRF, LSTM-CRF-ILP₁ produces a gain of 7.4 in event classification, 1.8 in event detection, and 4.6 in argument detection, with respect to the F1. We further investigate the effect of our heuristic method, LSTM-CRF-ILP_{multi} which deals with the multi-event sentence issue. Compared with other models, LSTM-CRF-ILP_{multi} selects several labeling sequences according to their objective value, and extract a number of events with comparable confidences from a sentence. As we can see from Table 3, this strategy may detect multi-type event mentions for a sentence, contributing to the increase of recalls, and F1 scores with a little drop of precisions.

4.5 Manual Evaluations

We randomly sample 150 unlabeled sentences from test data set. Annotators are asked to annotate the events and arguments to each sentence following two steps. First, determine whether a given sentence is positive or negative, and assign event types to positive sentences. Next, label all related arguments and their roles according to the types of events in the positive sentences. Each sentence is independently annotated by two annotators, and the inter-annotator agreement is 87% for event types and 79% for arguments.

Table 4 presents the average F1 score of manual evaluations. We can draw similar conclusions about the comparison of performances between different models as automatic evaluation. We demonstrate that LSTM-CRF-ILP_{multi} is the most effective model in event extraction as it achieves the highest F1 score in both manual and automatic evaluation.

S5: That night, in an apparent bid to kill Amos, the car instead runs over the sheriff, leaving Chief Deputy Wade Parent (played by James Brolin) in charge.			
Event Type	film.performance (Wrong labeled in data generation)		
Arguments	actor	character	film
	James Brolin	Wade Parent	the car
S6: Nicholas Hammond (born May 15, 1950) is an American actor and writer who is perhaps best known for his roles as Friedrich von Trapp in the film The Sound of Music, and as Peter Parker/Spider-Man on the CBS television series The Amazing Spider-Man.			
Event Type	film.performance		
Arguments	actor	character	film
	Nicholas Hammond	Friedrich von Trapp	The Sound of Music
Event Type	tv.regular_tv_appearance (Missing in generated data)		
Arguments	actor	character	series
	Nicholas Hammond	Peter Parker/Spider-Man	The Amazing Spider-Man

Figure 3: Example outputs of LSTM-CRF-ILP_{multi}.

Moreover, manual evaluation helps us to gain a deep insight of our data and models. We further conduct automatic evaluation on the manual annotated dataset and list the top 5 event types whose F1 scores of LSTM-CRF-ILP_{multi} differ greatly from automatic evaluation in Table 5.

Most of the performance differences are caused by the stage of data generation. Figure 3 examples two types of errors in data generation. Some of the sentences automatic generated test set are noisy, in other words, they do not imply any event while still match all key properties of certain instances. Take S5 as an example, though the phrases *the car* matches a film name, it does not indicate this film, and there is no explicit evidence expressing that an actor starring in a film. This is a bottleneck of our data generation strategy. During manual evaluation, we find 16 negative sentences which

¹<https://github.com/lzhang10/maxent>

Model	Event Classification			Argument Detection			Event Detection		
	P	R	F	P	R	F	P	R	F
CRF	96.8	9.93	18.0	64.8	6.54	11.9	29.8	3.06	5.55
MaxEnt	97.9	11.4	20.3	64.5	7.28	13.1	29.3	3.40	6.08
LSTM	97.2	62.4	75.1	77.1	53.9	63.5	51.0	32.8	39.9
LSTM-CRF	97.3	67.2	79.5	78.0	60.2	68.0	54.4	37.6	44.4
LSTM-CRF-ILP ₁	93.4	81.4	86.9	74.1	71.1	72.6	49.6	43.3	46.2
LSTM-CRF-ILP _{multi}	93.2	81.9	87.2	74.0	71.5	72.7	49.5	43.5	46.3

Table 3: Overall system performance of automatic evaluations. (%)

Model	EC	AD	ED
CRF	21.2	13.3	5.30
MaxEnt	17.7	11.7	5.44
LSTM	80.2	65.1	42.2
LSTM-CRF	81.6	68.6	44.1
LSTM-CRF-ILP ₁	85.4	70.2	44.2
LSTM-CRF-ILP _{multi}	85.5	70.4	44.6

Table 4: Average F1 scores of overall system performance of manual evaluations. (%) EC, AD, ED denote the event classification, argument detection and event detection, respectively.

are mistakenly labeled due to this reason. Unfortunately, our model fails to rectify 10 of them.

Remarkably, our LSTM-CRF-ILP_{multi} model can help find more CVT instances that not referenced in Freebase. There are two events mentioned in S6, while the arguments of the second event do not match any CVT instances in Freebase, leading to an omitting event in data generation. This phenomenon suggests that learning from distant supervision provided by Freebase, our model can help complete and update properties of Freebase instances in return.

Event type	P	R	F
olympics.medal_honor ²	↓ 25.0%	↓ 5.0%	↓ 13.8%
film.performance	↓ 21.4%	↑ 3.1%	↓ 10.3%
business.acquisition	→	↓ 7.1%	↓ 5.4%
tv.appearance ³	↓ 9.5%	↑ 3.0%	↓ 3.1%
film.release ⁴	↓ 7.7%	↑ 5.6%	↓ 0.55%

Table 5: Top 5 event types whose performances on event classification differ most from automatic evaluation. The model we evaluated is LSTM-CRF-ILP_{multi}

5 Related Work

Most event extraction works are within the tasks defined by several evaluation frameworks (e.g., MUC [Grishman and Sundheim, 1996], ACE [Doddington *et al.*, 2004], ERE [Song *et al.*, 2015] and TAC-KBP [Mitamura *et al.*, 2015]), all of which can be considered as a template-filling-based extraction task. These frameworks focus on limited number

of event types, which are designed and annotated by human experts and hard to generalize to other domains. Furthermore, existing extraction systems, which usually adopt a supervised learning paradigm, have to rely on those high-quality training data within those frameworks, thus hard to move to more domains in practice, regardless of feature-based [Gupta and Ji, 2009; Hong *et al.*, 2011; Li *et al.*, 2013] or neural-network-based methods [Chen *et al.*, 2015; Nguyen *et al.*, 2016].

Besides the works focusing on small human-labeled corpus, Huang *et al.* [2016] propose a novel Liberal Event Extraction paradigm which automatically discovers event schemas and extract events simultaneously from any unlabeled corpus. In contrast, we propose to exploit existing structured knowledge bases, e.g., Freebase, to automatically discover types of events as well as their corresponding argument settings, without expert annotation, and further automatically construct training data, with the essence of distant supervision [Mintz *et al.*, 2009].

Distant supervision (DS) has been widely used binary relation extraction, where the key assumption is that sentences containing both the subject and object of a $\langle subj, rel, obj \rangle$ triple can be seen as its support, and further used to train a classifier to identify the relation rel . However, this assumption does not fit to our event extraction scenario, where an event usually involves several arguments and it is hard to collect enough training sentences with all arguments appearing in, as indicated by the low coverage of **H1**. We therefore investigate different hypotheses for event extraction within the DS paradigm and propose to utilize time and syntactic clues to refine the DS assumption for better data quality. We further relieve the reliance on event trigger annotations by previous event extractors, and define a novel event extraction paradigm with key arguments to characterize an event type.

6 Conclusions and Future Work

In this paper, we propose a new event extraction paradigm without expert-designed event templates by leveraging structured knowledge bases to automatically acquire event schema and corresponding training data. We propose an LSTM-CRF model with ILP-based post inference to extract events without explicit trigger annotations. Experimental results on both manual and automatic evaluations show that it is possible to learn to extract KB-style events on automatically constructed training data. Furthermore, our model can extract information not covered by Freebase which indicates the possibility to extend this work to knowledge base population.

²The full name is olympics.olympic_medal_honor in Freebase.

³The full name is tv.regular_tv_appearance in Freebase.

⁴The full name is film.film_regional_release_date in Freebase.

References

- [Ahn, 2006] David Ahn. The stages of event extraction. In *Proceedings of the Workshop on Annotating and Reasoning about Time and Events*, pages 1–8. Association for Computational Linguistics, 2006.
- [Berant *et al.*, 2013] Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on freebase from question-answer pairs. In *EMNLP 2013*, volume 2, page 6, 2013.
- [Berger *et al.*, 1996] Adam L Berger, Vincent J Della Pietra, and Stephen A Della Pietra. A maximum entropy approach to natural language processing. *Computational linguistics*, 22(1):39–71, 1996.
- [Bollacker *et al.*, 2008] Kurt Bollacker, Colin Evans, Praveen Paritosh, Tim Sturge, and Jamie Taylor. Freebase: a collaboratively created graph database for structuring human knowledge. In *SIGMOD 2008*, pages 1247–1250. ACM, 2008.
- [Chen *et al.*, 2015] Yubo Chen, Liheng Xu, Kang Liu, Daojian Zeng, and Jun Zhao. Event extraction via dynamic multi-pooling convolutional neural networks. In *ACL 2015*, pages 167–176, 2015.
- [Collobert *et al.*, 2011] Ronan Collobert, Jason Weston, Léon Bottou, Michael Karlen, Koray Kavukcuoglu, and Pavel Kuksa. Natural language processing (almost) from scratch. *Journal of Machine Learning Research*, 12(Aug):2493–2537, 2011.
- [Doddington *et al.*, 2004] George R Doddington, Alexis Mitchell, Mark A Przybocki, Lance A Ramshaw, Stephanie Strassel, and Ralph M Weischedel. The automatic content extraction (ace) program-tasks, data, and evaluation. In *LREC 2004*, volume 2, page 1, 2004.
- [Grishman and Sundheim, 1996] Ralph Grishman and Beth Sundheim. Message understanding conference-6: A brief history. In *COLING 1996*, volume 96, pages 466–471, 1996.
- [Gupta and Ji, 2009] Prashant Gupta and Heng Ji. Predicting unknown time arguments based on cross-event propagation. In *ACL-IJCNLP 2009*, pages 369–372. Association for Computational Linguistics, 2009.
- [Gurobi Optimization, 2016] Inc. Gurobi Optimization. Gurobi optimizer reference manual, 2016.
- [Hochreiter and Schmidhuber, 1997] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- [Hong *et al.*, 2011] Yu Hong, Jianfeng Zhang, Bin Ma, Jianmin Yao, Guodong Zhou, and Qiaoming Zhu. Using cross-entity inference to improve event extraction. In *ACL 2011*, pages 1127–1136. Association for Computational Linguistics, 2011.
- [Huang *et al.*, 2015] Zhiheng Huang, Wei Xu, and Kai Yu. Bidirectional lstm-crf models for sequence tagging. *arXiv preprint arXiv:1508.01991*, 2015.
- [Huang *et al.*, 2016] Lifu Huang, T Cassidy, X Feng, H Ji, CR Voss, J Han, and A Sil. Liberal event extraction and event schema induction. In *ACL 2016*, 2016.
- [Kudo, 2005] Taku Kudo. Crf++: Yet another crf toolkit. *Software available at <http://crfpp.sourceforge.net>*, 2005.
- [Lafferty *et al.*, 2001] John Lafferty, Andrew McCallum, and Fernando Pereira. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *ICML 2001*, volume 1, pages 282–289, 2001.
- [Lample *et al.*, 2016] Guillaume Lample, Miguel Ballesteros, Sandeep Subramanian, Kazuya Kawakami, and Chris Dyer. Neural architectures for named entity recognition. *arXiv preprint arXiv:1603.01360*, 2016.
- [Li *et al.*, 2013] Qi Li, Heng Ji, and Liang Huang. Joint event extraction via structured prediction with global features. In *ACL 2013*, pages 73–82, 2013.
- [Manning *et al.*, 2014] Christopher D Manning, Mihai Surdeanu, John Bauer, Jenny Rose Finkel, Steven Bethard, and David McClosky. The stanford corenlp natural language processing toolkit. In *ACL 2014*, pages 55–60, 2014.
- [Mikolov *et al.*, 2013] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *NIPS 2013*, pages 3111–3119, 2013.
- [Mintz *et al.*, 2009] Mike Mintz, Steven Bills, Rion Snow, and Dan Jurafsky. Distant supervision for relation extraction without labeled data. In *ACL-IJCNLP 2009*, pages 1003–1011. Association for Computational Linguistics, 2009.
- [Mitamura *et al.*, 2015] Teruko Mitamura, Yukari Yamakawa, Susan Holm, Zhiyi Song, Ann Bies, Seth Kulick, and Stephanie Strassel. Event nugget annotation: Processes and issues. In *Proc. of the Workshop on EVENTS at the NAACL-HLT*, pages 66–76, 2015.
- [Nguyen *et al.*, 2016] Thien Huu Nguyen, Kyunghyun Cho, and Ralph Grishman. Joint event extraction via recurrent neural networks. In *NAACL-HLT 2016*, pages 300–309, 2016.
- [Rabiner, 1989] Lawrence R Rabiner. A tutorial on hidden markov models and selected applications in speech recognition. *Proc. of the IEEE*, 77(2):257–286, 1989.
- [Song *et al.*, 2015] Zhiyi Song, Ann Bies, Stephanie Strassel, Tom Riese, Justin Mott, Joe Ellis, Jonathan Wright, Seth Kulick, Neville Ryant, and Xiaoyi Ma. From light to rich ere: annotation of entities, relations, and events. In *Proc. of the Workshop on EVENTS at the NAACL-HLT*, pages 89–98, 2015.
- [Zeng *et al.*, 2015] Daojian Zeng, Kang Liu, Yubo Chen, and Jun Zhao. Distant supervision for relation extraction via piecewise convolutional neural networks. In *EMNLP 2015*, pages 1753–1762, 2015.